
Absence Is Not Omission: Availability, Vintage, and the Measurement of Selective Benchmark Disclosure

Anonymous Author(s)
Affiliation
Address
email

Abstract

1 Providers choose which benchmark results a release document prints; an
2 independent evaluator scores the same models on many more. Dating the
3 instruments seems to answer whether omitted results stand below printed
4 ones: a benchmark published after a model shipped could not be reported.
5 The comparison fails as a check on itself. Benchmarks postdating a release
6 already trail those predating it by 12.23 percentile points before any la-
7 bel enters, because the date deciding availability also fixes where a model
8 sits in the benchmark’s comparison pool, an age-period-cohort collinearity
9 reweighting cannot remove. The failure appears on a leaderboard we did
10 not build: across fourteen HELM Lite releases, fixed models keep identi-
11 cal scores while every win rate moves; a pool-invariant companion moves
12 nothing. Hand-coding every locatable release document, with an indepen-
13 dent second coding, shows the consequence: the coded gap rejects zero
14 at +11.4, yet with labels removed the same releases return +14.1, the dif-
15 ference exactly what composition implies. A study built this way turns
16 pool composition into apparent concealment. Reporting is performance-
17 related, reported benchmarks standing 8.7 points above omitted eligible
18 ones within release, yet no possession-grounded test detects withholding,
19 the change-based gap keeping the predicted sign inside its null. Carrying
20 a fixed suite across releases would make later silence informative; adding
21 labels to a moving pool does not.

22 1 Introduction

23 For each frontier model release, providers disclose benchmark scores through a model card,
24 while third parties execute standardised benchmarks for the same models [Liang et al.,
25 2023, Epoch AI, 2026]. This paper asks whether benchmark availability can identify the
26 older economic question of whether the provider’s undisclosed scores are typically worse than
27 their reported scores. Only when the signal is common knowledge does disclosure unravel
28 [Grossman, 1981, Milgrom, 1981]. As Dye [1985] details, when receivers cannot distinguish
29 between senders who have no information and senders who have bad information, receivers
30 treat senders as the same, encouraging withholding. This indistinguishability extends to
31 any research that attempts to analyse what was withheld, since we rely on the endowment,
32 the signal a sender had at the time of writing, which we do not observe in the work below.

33 The presence of benchmark release dates gives a rare ability to bound the endowment
34 exogenously: when benchmarks are not shared before publication, any benchmark released
35 after the model release is necessarily not in the provider’s endowment, regardless of what

36 providers ran themselves. This lets us distinguish what was able to be disclosed from what
37 was unable to be disclosed.

38 First, we provide a new longitudinal study of benchmark disclosure that uses this bound to
39 construct a panel: from a fixed snapshot of an independent benchmarking hub containing 819
40 model-versions, we aggregate over scaffold variants to 353 provider releases, 46 organisations,
41 74 benchmarks, and 3,082 model-benchmark scores. We provide a verified release date
42 for each benchmark, creating a panel of 2,355 eligible model-benchmark scores and 476
43 postdating scores that could not have been disclosed.

44 Second, we show that while the ability to bound the endowment is a unique and powerful
45 feature to understand withholding, constructing a relative ranking of performance is required
46 to judge withholding. Our central analysis compares reported versus omitted benchmarks
47 that were able to be reported, and requires the relative ranking to be unbiased across the
48 two sides of availability. Instead, there is a 12.23 percentile point difference on average
49 between benchmarks created before the model release and those created after it, with pre-
50 release benchmarks outperforming post-release benchmarks in 78.1% of the 146 releases that
51 have cells in both groups. None of this calculation uses whether the provider reported the
52 benchmark, and under the least biased ranking we can construct, the difference remains
53 at 9.14 points. This remains a structural problem, as the verification of the endowment
54 uses the same signal, the release date, that biases the relative ranking [Heckman, 1979].
55 In economic terms, benchmark vintages are a free signal to verify the endowment, which
56 Dye [1985] demonstrates cannot otherwise be known to receivers. However, the signal that
57 verifies the endowment is not excludable from the relative ranking, and is therefore self-
58 contaminating. This raises a fundamental design challenge for any research that verifies the
59 presence of information via vintage but prices evidence using a relative ranking, win rate,
60 or percentile. Not only must reporting decisions be observable, but also the relative ranking
61 must maintain a consistent reference group that is invariant to the endowment.

62 Third, we take the question outside our own panel. Across 14 releases of HELM Lite, the
63 pool grows from 30 models to 91 while 24 models keep bit-identical scores, and every one of
64 their published headline values moves, five pairs reversing order. A companion leaderboard
65 from the same team, behind an absolute headline, shows none of this.

66 Fourth, we code the record. Hand-coding the 100 releases with a locatable release-time
67 document, independently replicated on a fifth at $\kappa = 0.95$ on reported versus not, yields the
68 estimator such a study would use. Omitted benchmarks stand +11.4 points above those a
69 provider could not have reported, excluding zero yet short of its matched counterpart by
70 what composition implies: it cannot separate concealment from the availability gap. What
71 survives is composition: reported benchmarks stand 8.7 points above omitted eligible ones
72 within release ($p = 0.0001$), robust to benchmark fixed effects.

73 1.1 Related literature

74 That reporting in model documentation is imperfect, incomplete, and dependent on the
75 sender is established [Staufer et al., 2025, Kuehnert et al., 2026], with two papers focusing
76 on the reporting gaps themselves: Zhang et al. [2026] finds a benchmark family missing
77 from all main result tables across four flagship releases, and Ghosh et al. [2026] defines the
78 reporting standard that would make the omission observable. Both lack benchmark dates,
79 so there is no way to distinguish a skipped benchmark from one the release predates. Bean
80 et al. [2025] instead evaluates 445 benchmarks for construct validity, a different question
81 about whether instruments measure what they claim, while we take the instruments as
82 given.

83 Evidence is stronger on leaderboards: Singh et al. [2025] characterises selection in real time,
84 providers testing private variants and publishing the best, while Xu et al. [2025] shifted
85 HELM’s primary metric off a win rate that depends on the models compared. In both cases
86 the leaderboard changes for reasons other than the evidence, but neither isolates, and we
87 do at the cell level, an omission from a result that could not have existed. What no paper
88 here controls for is availability, which we introduce: absent a confirmed benchmark date,
89 the distinction between “was not reported” and “did not exist” collapses.

90 The economics literature frames the problem and gives a cautionary note: silence unravels
 91 when endowments are commonly observed [Grossman, 1981, Milgrom, 1981] but persists if
 92 the recipient cannot distinguish an empty endowment from a concealed one [Dye, 1985, Jung
 93 and Kwon, 1988] or when revelation incurs a cost [Verrecchia, 1983], frictions our design does
 94 not separate. Empirically, the warning bears out: unravelling does not occur in practice
 95 [Dranove and Jin, 2010, Jin and Leslie, 2003, Hotz and Xiao, 2013, Bederson et al., 2018]
 96 or in the laboratory [Jin et al., 2021]. What all this work misses is the endowment directly,
 97 which our bound aims to recover. Concurrent theory takes a direct approach to strategic
 98 assessment and persuasion [Truong et al., 2026, Dai et al., 2026]; that work is theory, not
 99 measurement.

100 2 Data and Setting

101 2.1 The disclosure problem

102 A release i ships at date t_i carrying a single document. Let $\mathcal{B}$ index benchmarks, each with
 103 its own publication date τ_b , let $E_i \subseteq \mathcal{B}$ be the endowment, the benchmarks whose result the
 104 provider holds when it writes, and let $D_i \subseteq E_i$ be the disclosed set, which reports a score
 105 s_{ib} for each $b \in D_i$. The receiver observes D_i and its scores and must split the complement
 106 into results held back and results never held.

107 Under a commonly known endowment, silence is priced at the worst value it could conceal
 108 and disclosure unravels to $D_i = E_i$ [Grossman, 1981, Milgrom, 1981]. Two distinct frictions
 109 break that. In Verrecchia [1983] the endowment is known and disclosure is costly, so with-
 110 holding is a cutoff in s_{ib} . In Section 3 of Dye [1985] disclosure is free but the receiver cannot
 111 tell an empty endowment from a concealed one, because a sender can prove possession and
 112 cannot prove emptiness.

113 Here an independent evaluator publishes s_{ib} for cells the provider passed over, and each τ_b
 114 makes infeasibility verifiable. Writing $A_i = \{b : \tau_b < t_i\}$ for the available set and O_i for the
 115 cells the evaluator scores,

$$D_i \subseteq E_i \subseteq A_i, \quad \theta = \Pr(b \notin D_i \mid b \in O_i \cap A_i).$$

116 A provider may hold results the evaluator never produced and a benchmark can exist and
 117 go unrun, so θ is defined against the evaluator’s coverage rather than the true endowment,
 118 and the Dye friction survives inside A_i .

119 Judging what was withheld then requires a measure of standing P_{ib} , since raw scores are not
 120 comparable across benchmarks. The design rests on an exclusion restriction: conditional on
 121 the innocent explanations, membership in A_i does not enter the equation for P_{ib} [Heckman,
 122 1979]. A date rule cannot satisfy it: $t_i - \tau_b$, t_i and τ_b are exactly collinear, so if any two
 123 of the three move P_{ib} the third cannot be excluded from it. All three move it here, which
 124 Section 3 shows and Appendix B bounds. Two things follow, and they are not the same. The
 125 first is practical: estimating θ needs disclosure labels, which the coded record of Section 5
 126 supplies. The second is structural: reading θ as evidence of withholding runs into A_i itself.
 127 Availability is a deterministic threshold in benchmark maturity, $\mathbf{1}[t_i - \tau_b > 0]$, so no cell of the
 128 same maturity sits on both sides of the boundary. Wherever standing depends on maturity,
 129 conditioning finely enough to make availability excludable destroys positivity. So availability
 130 alone cannot globally identify selective disclosure from any pool-relative standing measure,
 131 win rate or percentile alike, and a jointly estimated scale inherits the same selection through
 132 which cells exist. What survives, under continuity, is the discontinuity at the threshold.

133 2.2 The independent score set

134 We work with a pinned version of Epoch’s benchmarking hub scraped on August 17th, 2026.
 135 This data includes 819 model-versions. Since we are interested in release-level disclosure, we
 136 use organisation, model and release date to identify releases. Epoch uses a version attribute
 137 to distinguish different scaffolds within the same release, e.g. different reasoning effort or
 138 context length variants. Since the provider publishes one document for one release, we col-
 139 lapse these versions to releases, taking the maximum score among scaffolds and retaining the

spread as a check. We end up with 353 releases, 74 benchmarks and 3,082 release-benchmark cells that have an associated independent score. The data is clustered by provider.

2.3 Eligibility and benchmark dates

A provider cannot omit a benchmark that did not yet exist. To ensure temporal eligibility, we collect release dates for benchmarks: the earliest verifiable public release date (from Epoch if available, from the primary source of the benchmark if not, translating between preprint dates, code repository dates and different versions to get the first public release date). A cell is eligible only if the benchmark release date strictly predates the release’s date and if the benchmark’s release date is verified. This restriction assumes providers could not have access to unreleased benchmarks; if a provider was given early access to a benchmark, this assumption would not hold. Such violations shrink rather than create the reporting gap; they move available cells into the postdating category. Such violations leave no trace in our coded data: no document reports a benchmark before its public release date. This restriction leaves 2,355 out of the 3,082 cells eligible. Cells whose benchmark strictly postdates the release’s date form a separate set of 476, cells that could not have been reported. We drop cells without a verified benchmark date rather than imputing a date; inaccuracies in dates shrink the boundary estimate we construct in Section 3. For comparing releases, we do not require a minimum number of benchmarks; this would exclude the releases that only have one eligible and one postdating cell. Of the 146 releases with both eligible and postdating cells, 57 have fewer than eight scored benchmarks. The simpler obstacle, that how many benchmarks a release could report is set by when it shipped, is discussed in Appendix A.

To code disclosure, we adapt ORBIT, Outcome Reporting Bias In Trials. ORBIT is standardised in clinical trials, where disappearing outcomes (outcomes specified in the protocol but not reported in the publication) have a long literature on measurement [Chan et al., 2004, Kirkham et al., 2010, Dwan et al., 2013]. For existing cells, we record whether the provider reports a number, cites the benchmark without reporting, or reports a variant; for missing cells, we record whether a previous member of the same family reported it, whether contemporaries reported it, or whether no similar provider reported it. Of the 108 releases in the queue, 103 with eligible cells in our pinned dataset and five in a more recent vintage, we code the 100 for which we can locate a document from release time; 1,280 cells in total. None of the analysis before Section 5 depends on coded disclosure.

2.4 Measuring standing

Raw scores are not comparable across benchmarks, so standing is measured as a within-benchmark percentile computed against contemporaneous models, meaning those whose release dates fall within 182 days on either side of the focal release, with midranks for ties. The peer count and the share of the window released after the focal model are retained alongside the percentile. It is the measure a windowed comparison suggests, and the one the paper shows is biased in a way that mimics selective disclosure.

2.5 Inference

Inference is constrained by the same structure. The panel has 46 nominal provider clusters, but 24 contribute fewer than ten cells, 20 appear in a single release, and the five largest hold 68.1% of the sample. Cluster-robust asymptotics are not credible below twelve clusters, and we report no analytic standard error there [Cameron et al., 2008]. Regression coefficients in that range use the restricted wild cluster bootstrap over 9,999 draws, which imposes the null before resampling Rademacher weights at the provider level. Release-level means use a cluster sign-flip randomization test over 9,999 draws, reversing an organisation’s entire contribution at once; it preserves provider dependence but not benchmark dependence, which the two-way clustered regressions carry, and their intervals are percentile bootstraps over 10,000 provider resamples. Both report $(k + 1)/(\text{draws} + 1)$, with k the draws at least as extreme as the observed statistic. Mirrored sign vectors leave 2^{G-1} distinct statistics over G clusters, so a resolution of 0.0001 needs fifteen clusters; the 18 organisations behind the headline estimate clear it. Below twelve clusters the bootstrap switches to a six-point weight distribution whose finer support restores it [MacKinnon and Webb, 2016, 2018].

3 The Availability Gap

3.1 The placebo comparison

Separating results held back from results never held calls for a control group whose endowment is known to be empty [Dye, 1985]. The dates appear to supply one: a benchmark not yet published when the model shipped could not, absent prepublication access, have been reported, while an eligible one could have been. Whether it was reported is what a disclosure label would say, and none enters here. Whatever makes an eligible benchmark look bad for innocent reasons, difficulty drift, saturation, or a shift in what the field measures, presses on both sets alike, so under those explanations the difference should be zero. One innocent channel is asymmetric by construction: a benchmark published before a model shipped can sit in its training corpus and a later one cannot, so contamination raises eligible standing only, beyond any reweighting of the comparison.

It is not zero. Across 146 releases with both kinds of cell, eligible benchmarks outscore postdating ones by 12.23 percentile points on average, with a median of 10.69, and the gap is positive in 78.1% of releases. Nothing in that calculation knows which benchmarks were reported. A measure that balances the two sides of the window, described below, cuts the figure to 9.14 points, positive in 70.9% of 141 releases, without removing it. A cluster sign-flip randomization test over the 18 organisations holding both kinds of cell puts that mean at $p = 0.0001$, the smallest value the draws can return. Cells that could not have been reported already stand lower than cells that could, so the group is not a control.

3.2 Decomposing the gap

Table 1 locates the failure, regressing standing on a placebo indicator, equal to one when the benchmark postdates the release, always within release. Release effects alone leave -13.5 points. Benchmark fixed effects absorb 63.4% of that, and structurally so: placebo cells concentrate by construction in benchmarks that entered late, and late benchmarks are less saturated, with 8.3% of scores on the later half of the 54 benchmarks with ten or more scores in the top decile of the benchmark’s own range against 13.2% on the earlier half ($p = 0.015$), more room left below the ceiling. The peer-window asymmetry term, the share of a cell’s comparison window released after the focal model, absorbs 39.8% on its own; conditioning jointly on benchmark effects and both window terms absorbs all but 0.003 points, an interval reaching $+3.4$, a residual bounded under a third of the original gap rather than at zero.

Conditioning set	Placebo coefficient	t	Absorbed
Release fixed effects only	-13.529 (2.038)	-6.637	0.0%
Release FE + benchmark fixed effects	-4.949 (2.183)	-2.267	63.4%
Release FE + peer-window asymmetry	-8.151 (1.865)	-4.370	39.8%
Release FE + peer count	-2.107 (1.882)	-1.120	84.4%
Two-way fixed effects + both window terms	-0.003 (1.748)	-0.002	100.0%

Table 1: Placebo coefficient by conditioning set, percentile points, within release. Two-way provider-by-benchmark standard errors (45 by 61 clusters; provider-only in Appendix B), $n = 2,831$ cells. Absorbed is the fall in absolute value from the first row; peer count alone is a suppressor, so the last row conditions on both window terms jointly.

The mechanism behind the second dimension is geometric. The comparison window is symmetric in calendar time, 182 days on each side of the release, but a benchmark’s model coverage is not symmetric in calendar time, because coverage begins when the benchmark is built. Figure 1 plots the two objects this creates against benchmark maturity at release: the share of the window that is newer falls with maturity and jumps at the availability boundary, while standing runs through it without a detectable break. The mechanism is discontinuous where the outcome is not; peer-set composition is a property of timing alone.

The share of a cell’s window released after the focal model is 0.4683 for eligible cells and 0.7065 for placebo cells, against 0.5 for a two-sided window. Within release, the slope of standing on that share is -30.63 percentile points per unit among eligible cells alone. Figure 2 plots standing against that share for both groups at once. Eligible and placebo cells do not form two clouds at different heights; they trace a single declining curve, common only




Figure 1: The mechanism and the outcome at the availability boundary. Top: the share of a cell’s comparison window released after the focal model, against benchmark maturity at release; it falls with maturity and jumps at the boundary. Bottom: within-benchmark standing on the same axis, continuous through it. Sixty-day bins; intervals clustered on organisation.

once benchmark composition is absorbed (Table 1), and differ in where along it they fall. On the variable that assigns them, benchmark maturity at release, they cannot overlap, so the contrast is an extrapolation across that boundary rather than a comparison within common support. The check is at the boundary, the one place Section 2.1 leaves identified under continuity [Hahn et al., 2001]. Nothing detectable happens there. Local linear estimates are negative at every bandwidth and order, every interval covers zero though the widest reaches $+8.05$, and a window chosen on covariate balance puts the difference at -2.52 , randomization p 0.447 [Cattaneo et al., 2015]. The two groups’ mean maturity differs by 793 days; the gap lives in that distance, not at the boundary (balance tests in Appendix B).




Figure 2: Standing against the share of a cell’s comparison window released after the focal model; dashed verticals mark each group’s mean share, and the black dashed line the symmetric 0.5. Eligible and placebo cells lie on one declining curve, so the contrast between them is a position on that curve rather than a fact about disclosure. Intervals are clustered on organisation.

Standing decomposes into a weighted average of position among older and newer peers, $\text{percentile} = (1 - w) P_{\text{old}} + w P_{\text{new}}$ with w the newer share of window peers, exact cell by cell. Replacing the empirical weight with one half defines the side-balanced measure, available for 92.5% of all cells and 92.8% of eligible-or-placebo cells, and cuts the same slope to -8.56 .

That the identity holds does not make the gradient an accounting tautology. Shuffling scores within benchmark destroys the capability trend while leaving dates, window membership, peer counts and newer shares bit-identical. Under that null, the slope has mean -0.09 and standard deviation 3.12. The observed -30.63 sits 9.8 standard deviations away, and none of 300 draws reaches it. Replacing scores instead with a within-benchmark linear fit on date plus permuted residuals, so that a capability trend exists and nothing else does, recovers 93% of the observed slope, averaged over 200 draws. Comparison-set geometry is inert on its own and becomes a bias only when it multiplies a real trend in capability.

3.3 Trimming and reweighting

The phenomenon has a name in nonparametric estimation: local averages are biased near the edge of a support, and the standard remedies recentre rather than discard [Gasser and Müller, 1979, Fan, 1992]. What is unusual is the boundary’s coincidence with the

endowment: a benchmark was available to a release exactly when the model did not sit at the left edge of that benchmark’s coverage. The indicator of what a provider could have held is therefore the indicator of where the window is one-sided, and removing the edge removes the contrast. The coincidence has an older name, the age-period-cohort identity: release date, benchmark vintage and their difference are exactly collinear, and each has a measured signature here: t_i in the capability trend, τ_b in saturation, $t_i - \tau_b$ in window composition. Only the linear components are unidentified in that structure [Fosse and Winship, 2019a,b], so what fails is a slope. Three sign restrictions bound it: capability rises with the calendar, later benchmarks are less saturated, and standing rises with maturity. All three are assumptions on unidentified components, each motivated by a measurement of an identified quantity. Under them the unidentified component accounts for at most 24 to 34% of the placebo gap across specifications, roughly three percentile points, which is where the saturated regression also lands. The remainder is carried by the identified reduced-form slopes, and Appendix B gives the derivation.

Standing measure	Mean	Median	Positive	Sign-flip p	Releases	Cells
Windowed percentile, $\pm 182d$	+12.23	+10.69	78.1%	0.0001	146	100.0%
Rank against all models ever	+22.75	+22.66	89.0%	0.0001	146	100.0%
Trim to two-sided windows	+9.96	+10.69	72.2%	0.0001	115	70.4%
Side-balanced percentile	+9.14	+9.62	70.9%	0.0001	141	92.8%

Table 2: The label-free placebo null under each candidate measure: release-level mean standing of eligible minus postdating benchmarks, in percentile points. Every entry should be zero. Sign-flip p is the provider-clustered test of Section 2.5; Cells, the share of cells carrying the measure.

Table 2 shows what follows. Ranking against every model ever scored nearly doubles the contamination; trimming to two-sided windows discards 30% of the cells to remove 19%, its row the uncorrected gap on the retained cells, composition rather than correction. Shortening the window does reduce the gap, to +8.59 at 91 days, but in proportion to the asymmetry it removes rather than by breaking the mechanism. Across four widths and four benchmark-count minima the gap is positive in all 16 specifications, from +8.59 to +21.66, and tracks the eligible-minus-placebo gap in window composition at $r = 0.991$. Dropping any one of the 18 organisations leaves it between +11.45 and +13.20. The correction is a covariate, not a subsample.

One reading of that number would be too generous to it. Window composition is a function of the same dates that define the placebo indicator, so conditioning on it does not adjust for a confounder sitting outside the comparison; it splits the gap into the channel carrying it. A study that holds disclosure labels and wants an estimate of θ must first argue that window composition is settled before a provider decides what to report, or it will condition away part of what it is measuring. Provider identity explains little of the variance in window composition; Appendix B reports the decomposition and the coverage tilt it does not correct.

The side-balanced percentile is the one change of measure doing substantial work, cutting the placebo mean to +9.14 and the eligible-cell gradient to -8.56 . Figure 4 in Appendix B puts it against the uncorrected measure, near flat where cells are dense. Reweighting cannot reach what remains: balancing repairs the weight on each side of a model but not which models an evaluator chose to score early, so evaluator attention stays a confound. Undefined wherever a side is empty, it accompanies rather than replaces the uncorrected measure.

4 Evidence from HELM

Everything so far concerns a percentile constructed here on a panel assembled here, so the effect may belong to the construction rather than to leaderboards in general. HELM Lite publishes its full accuracy table at every versioned release, archives the earlier ones, and reports a pool-relative headline [Liang et al., 2023], so the same evidence can be read at 14 points in time while the pool grows underneath it.

If a leaderboard moves, the first thing to rule out is a revised suite. Across releases v1.0.0 through v1.13.0, the ten scenario columns are identical, with nothing added, dropped or renamed beyond one relabelled column header. The same 24 models appear in all 14 releases

306 with their 240 published scenario scores bit-identical throughout, while the evaluated pool
307 grows from 30 models to 91.

308 Against that frozen evidence, all 24 constant models see their published Mean win rate
309 change, by 0.1232 on average in absolute value and by as much as 0.2178 for Mixtral 8x7B.
310 Most of that movement is arithmetic: Mean win rate is the fraction of head-to-head compar-
311 isons a model wins, so a fixed model faces a larger, later field as the pool grows; the change
312 correlates with initial standing at -0.766 .

313 The movement that matters is elsewhere. If the drift were a monotone transform of the
314 frozen evidence, it would be harmless, since a monotone rescaling preserves every order. It
315 is not. Of the 24, 22 fall and two rise, and five of the 276 pairs among them, 1.8%, reverse
316 between the first release and the last, each a published order inversion with bit-identical
317 evidence beneath it. Figure 3, in Appendix B, plots each model’s rank among the constant
318 set at the first release against the last, so a reversed pair is a crossing, countable on the
319 page while the evidence beneath is fixed. Unlike test-set reuse [Roelofs et al., 2019], the
320 evidence is fixed; the reference set is what moves. The published numbers are correct as
321 defined at every release; what changed was the pool, and the published order followed it.
322 Singh et al. [2025] reports a comparable pattern from a different mechanism, deprecation
323 policy rather than a coverage boundary changing which models remain available, and in
324 both the reference set rather than the evidence moves the ranking.

325 The control, HELM Capabilities, comes from the same organisation, whose maintainers had
326 diagnosed the win rate as pool-dependent [Xu et al., 2025]; it runs on the same infrastructure
327 and cadence, its pool growing comparably, from 22 models to 68 over 16 releases. Its headline
328 is an absolute mean over scenarios rather than a fraction of comparisons won. The same 22
329 models appear in every release with 110 scenario scores bit-identical, none of the 22 headline
330 values moves, and no pair reorders, at the endpoints or ever. An absolute mean over bit-
331 identical scores on a fixed suite cannot move with the pool, so zero drift is arithmetic, a
332 scope condition rather than a test that could have failed.

333 This need not be read as an error on HELM’s part: the absolute headline on its own
334 Capabilities leaderboard is the point. A pool-relative headline is a joint statement about
335 a model and about the population evaluated beside it; read later as statements about the
336 models alone, two such numbers move standing between providers for reasons belonging to
337 neither. Restoring comparability across moments is the classical equating problem [Kolen
338 and Brennan, 2014], avoided by a scale estimated jointly [Ho et al., 2025]. Panel and
339 leaderboard fail identically: the reference set, not the evidence, moves the number.

340 5 Coding the Record

341 A benchmark can be missing from a model card because it was irrelevant, never run, or
342 run and withheld. Only the third is disclosure in the sense of Grossman [1981], Milgrom
343 [1981] and Dye [1985], and only it is a reader’s to price. The within-family drop narrows
344 it: a provider that reports b at t and omits it at $t + 1$, an independent score present,
345 has established relevance and possession of the harness, with the independent score as the
346 registry-style counterfactual [Chan et al., 2004, Dwan et al., 2013, Kirkham et al., 2010]. It
347 does not establish that b was run for $t + 1$, so the Dye [1985] uncertainty narrows rather
348 than closes; its ceiling, 37 pairs, and detection threshold, roughly 38 points at one drop, sit
349 in Appendix B.

350 A single coder read the release-time document of each of the 100 releases for which one is
351 locatable; eight of the queued 108 have none. The record notes every benchmark a document
352 reports, with variants, bare mentions, and stated benign absences; the ORBIT categories
353 are derived mechanically from these records and family history. This yields 1,280 coded
354 cells; 31 from five later-vintage releases sit outside the pinned panel and are excluded. A
355 second coder then recoded a provider-stratified fifth of the releases, blind to the first coding
356 and without calibration contact; Appendix B records the one release where that blindness
357 broke. Agreement is Cohen’s κ over the derived cells [Cohen, 1960], with intervals from
358 resampling releases. The two codings agree at 0.95 [0.80, 1.00] on the primitive judgment of
359 whether a benchmark was reported, at 0.85 [0.72, 0.95] over the full category set, and at 0.81

360 [0.48, 1.00] on the high-versus-low-suspicion split. The single E cell each coding derives is the
361 same cell, encouraging but no reliability estimate for so rare a category, and Appendix B
362 reports the confusion structure. Among the 1,249 eligible coded cells that resolve in the
363 panel, 76.9% report no number, and the release-level mean is 73.1% [65.9, 79.2]. The rate is
364 descriptive and is not, by itself, evidence of selection.

Standing measure	Coded gap	95% CI	No labels	Difference
Windowed percentile, $\pm 182d$	+11.40	[+5.6, +17.0]	+14.13	-2.73
Side-balanced percentile, $\pm 182d$	+7.90	[-0.4, +15.2]	+9.82	-1.92
Rank against all models ever	+22.33	[+18.0, +28.1]	+24.50	-2.17
Windowed percentile, $\pm 90d$	+8.15	[+1.3, +14.7]	+10.36	-2.21
Windowed percentile, $\pm 365d$	+18.03	[+13.9, +23.1]	+20.53	-2.50

Table 3: The coded omission deficit beside the label-free gap, in percentile points, over 71 releases and 13 providers (70 side-balanced). No labels is the same contrast with no disclosure coding, on the same releases; Difference is Coded gap minus No labels. Intervals are provider-clustered bootstraps; four of five coded gaps exclude zero.

365 Table 3 is the section’s result. The disclosed-versus-omitted comparison returns +8.7 points
366 [4.1, 13.9] over 63 releases; the intended availability contrast, omitted-eligible against post-
367 dating cells over the 71 releases holding both, returns +11.4 [5.6, 17.0] and would reject a
368 zero null at $p = 0.003$. Read against Section 3, it is the availability gap: the label-free
369 contrast on the same releases is +14.1, and under every measure and window the coded gap
370 falls 1.9 to 2.7 points short. The shortfall is an exact accounting identity, the average over
371 the same releases of the reported share times the within-release reported-omitted gap, +2.7,
372 so the labels change this contrast only through composition.

373 What survives is composition. Reported benchmarks stand 8.7 points above omitted eligi-
374 ble ones within release, a gap larger than all 9,999 count-preserving random assignments
375 ($p = 0.0001$, the draws’ floor; null standard deviation 1.9) and one the availability gap can-
376 not produce. It survives coding-error, variant-reclassification, and leave-one-out exercises
377 because the availability gap does; benchmark identity does not carry it, since benchmark
378 fixed effects beside the release effects raise the coefficient from +7.4 to +9.3 and the window-
379 share control leaves +7.4. It does not identify whether standing drove the reporting: rel-
380 evance, salience, and within-eligible contamination can move disclosure and performance
381 together.

382 Possession is the remaining distance; the original drop estimator and four possession cuts
383 return one verdict. The drop estimator realises 16 drops across 15 releases and returns
384 +2.9, opposite the withholding prediction, sign-flip $p = 0.25$ over five providers, the null
385 its calibration priced. No document states that a benchmark was evaluated while reporting
386 neither a score nor an outcome, across the 295 cells documenting a run. A change-based test,
387 each drop against its standing under the previous release, gives -3.7, the predicted sign,
388 at permutation $p = 0.16$ over 42 transitions; possession tiers and equal-sized substitutions
389 detect nothing further. Appendix B details these and three unrepaired limits: unlocatable
390 documents, coverage that may follow provider emphasis, and the best-across-scaffolds rule.

391 6 Conclusion

392 We set out to measure how much of what providers could report they choose not to. The
393 coded record returns the number and the reason to distrust it: +11.4 points, rejecting
394 zero yet short of its label-free counterpart by what composition implies, too imprecise to
395 exclude a strategic component of a few points. A study built this way turns composition
396 into apparent concealment. Three claims separate: non-reporting is extensive, composition
397 is performance-related, and withholding is not established, every possession-grounded test
398 returning null, absent, or opposite-signed evidence. Dating the instruments bounds the
399 endowment, an access assumption, and they fix where a model sits in a benchmark’s coverage:
400 what identifies the endowment moves the measure pricing silence. Growing the 16-drop
401 record is a reporting-standard question: a jointly estimated scale [Ho et al., 2025] shrinks
402 the placebo gap without closing it, and a fixed suite carried across releases makes later
403 silence informative, narrowing the empty-endowment excuse and Dye [1985] pooling.

References

- Andrew M. Bean, Ryan Othniel Kearns, Angelika Romanou, Franziska Sofia Hafner, Harry Mayne, et al. Measuring what matters: Construct validity in large language model benchmarks. In *Advances in Neural Information Processing Systems* 38, 2025. arXiv:2511.04703.
- Benjamin B. Bederson, Ginger Zhe Jin, Phillip Leslie, Alexander J. Quinn, and Ben Zou. Incomplete disclosure: Evidence of signaling and countersignaling. *American Economic Journal: Microeconomics*, 10(1):41–66, 2018. doi: 10.1257/mic.20150178.
- A. Colin Cameron, Jonah B. Gelbach, and Douglas L. Miller. Bootstrap-based improvements for inference with clustered errors. *Review of Economics and Statistics*, 90(3):414–427, 2008. doi: 10.1162/rest.90.3.414.
- Matias D. Cattaneo, Brigham R. Frandsen, and Rocío Titiunik. Randomization inference in the regression discontinuity design: An application to party advantages in the U.S. senate. *Journal of Causal Inference*, 3(1):1–24, 2015. doi: 10.1515/jci-2013-0010.
- An-Wen Chan, Asbjørn Hróbjartsson, Mette T. Haahr, Peter C. Gøtzsche, and Douglas G. Altman. Empirical evidence for selective reporting of outcomes in randomized trials. *JAMA*, 291(20):2457–2465, 2004. doi: 10.1001/jama.291.20.2457.
- Jacob Cohen. A coefficient of agreement for nominal scales. *Educational and Psychological Measurement*, 20(1):37–46, 1960.
- Yifan Dai, Drew Fudenberg, and Harry Pei. Bayesian persuasion with selective disclosure, 2026. arXiv:2601.05914.
- David Dranove and Ginger Zhe Jin. Quality disclosure and certification: Theory and practice. *Journal of Economic Literature*, 48(4):935–963, 2010. doi: 10.1257/jel.48.4.935.
- Kerry Dwan, Carrol Gamble, Paula R. Williamson, and Jamie J. Kirkham. Systematic review of the empirical evidence of study publication bias and outcome reporting bias. *PLoS ONE*, 8(7):e66844, 2013. doi: 10.1371/journal.pone.0066844.
- Ronald A. Dye. Disclosure of nonproprietary information. *Journal of Accounting Research*, 23(1):123–145, 1985. doi: 10.2307/2490910.
- Epoch AI. Ai benchmarking hub. <https://epoch.ai/benchmarks>, 2026. Data snapshot of 17 August 2026, CC-BY.
- Jianqing Fan. Design-adaptive nonparametric regression. *Journal of the American Statistical Association*, 87(420):998–1004, 1992. doi: 10.1080/01621459.1992.10476255.
- Ethan Fosse and Christopher Winship. Bounding analyses of age-period-cohort effects. *Demography*, 56(5):1975–2004, 2019a. doi: 10.1007/s13524-019-00801-6.
- Ethan Fosse and Christopher Winship. Analyzing age-period-cohort data: A review and critique. *Annual Review of Sociology*, 45:467–492, 2019b. doi: 10.1146/annurev-soc-073018-022616.
- Theo Gasser and Hans-Georg Müller. Kernel estimation of regression functions. In *Smoothing Techniques for Curve Estimation*, volume 757 of *Lecture Notes in Mathematics*, pages 23–68. Springer, 1979. doi: 10.1007/BFb0098489.
- Avijit Ghosh, Anka Reuel, Jenny Chim, Wm. Matthew Kennedy, et al. Evaluation cards: An interpretive layer for ai evaluation reporting, 2026. arXiv:2606.09809.
- Sanford J. Grossman. The informational role of warranties and private disclosure about product quality. *Journal of Law and Economics*, 24(3):461–483, 1981. doi: 10.1086/466995.
- Jinyong Hahn, Petra Todd, and Wilbert Van der Klaauw. Identification and estimation of treatment effects with a regression-discontinuity design. *Econometrica*, 69(1):201–209, 2001. doi: 10.1111/1468-0262.00183.

451 James J. Heckman. Sample selection bias as a specification error. *Econometrica*, 47(1):
452 153–161, 1979. doi: 10.2307/1912352.

453 Anson Ho, Jean-Stanislas Denain, David Atanasov, Samuel Albanie, and Rohin Shah. A
454 rosetta stone for AI benchmarks, 2025. arXiv:2512.00193.

455 V. Joseph Hotz and Mo Xiao. Strategic information disclosure: The case of multiattribute
456 products with heterogeneous consumers. *Economic Inquiry*, 51(1):865–881, 2013. doi:
457 10.1111/j.1465-7295.2010.00340.x.

458 Ginger Zhe Jin and Phillip Leslie. The effect of information on product quality: Evidence
459 from restaurant hygiene grade cards. *Quarterly Journal of Economics*, 118(2):409–451,
460 2003. doi: 10.1162/003355303321675428.

461 Ginger Zhe Jin, Michael Luca, and Daniel Martin. Is no news (perceived as) bad news?
462 an experimental investigation of information disclosure. *American Economic Journal:*
463 *Microeconomics*, 13(2):141–173, 2021. doi: 10.1257/mic.20180217.

464 Woon-Oh Jung and Young K. Kwon. Disclosure when the market is unsure of information
465 endowment of managers. *Journal of Accounting Research*, 26(1):146–153, 1988.

466 Jamie J. Kirkham, Kerry M. Dwan, Douglas G. Altman, Carrol Gamble, Susanna Dodd,
467 Rebecca Smyth, and Paula R. Williamson. The impact of outcome reporting bias in
468 randomised controlled trials on a cohort of systematic reviews. *BMJ*, 340:c365, 2010. doi:
469 10.1136/bmj.c365.

470 Michael J. Kolen and Robert L. Brennan. *Test Equating, Scaling, and Linking: Methods*
471 *and Practices*. Springer, 3rd edition, 2014. doi: 10.1007/978-1-4939-0317-7.

472 Blaine Kuehnert, Nari Johnson, Ravit Dotan, and Hoda Heidari. Disclosure or mar-
473 keting? analyzing the efficacy of vendor self-reports for vetting public-sector ai. In
474 ACM Conference on Fairness, Accountability, and Transparency (FAccT), 2026. doi:
475 10.1145/3805689.3812220.

476 Percy Liang, Rishi Bommasani, Tony Lee, et al. Holistic evaluation of language models.
477 *Transactions on Machine Learning Research*, 2023. arXiv:2211.09110.

478 Wei Ma and George H. Chen. Missing not at random in matrix completion: The effectiveness
479 of estimating missingness probabilities under a low nuclear norm assumption. In *Advances*
480 *in Neural Information Processing Systems 32*, 2019. arXiv:1910.12774.

481 James G. MacKinnon and Matthew D. Webb. Wild bootstrap inference for wildly different
482 cluster sizes. *Journal of Applied Econometrics*, 32(2):233–254, 2016. doi: 10.1002/jae.
483 2508.

484 James G. MacKinnon and Matthew D. Webb. The wild bootstrap for few (treated) clusters.
485 *The Econometrics Journal*, 21(2):114–135, 2018. doi: 10.1111/ectj.12107.

486 Justin McCrary. Manipulation of the running variable in the regression discontinuity design:
487 A density test. *Journal of Econometrics*, 142(2):698–714, 2008. doi: 10.1016/j.jeconom.
488 2007.05.005.

489 Paul R. Milgrom. Good news and bad news: Representation theorems and applications.
490 *Bell Journal of Economics*, 12(2):380–391, 1981. doi: 10.2307/3003562.

491 Rebecca Roelofs, Vaishaal Shankar, Benjamin Recht, Sara Fridovich-Keil, Moritz Hardt,
492 John Miller, and Ludwig Schmidt. A meta-analysis of overfitting in machine learning. In
493 *Advances in Neural Information Processing Systems 32*, 2019.

494 Shivalika Singh, Yiyang Nan, Alex Wang, Daniel D’Souza, Sayash Kapoor, Ahmet Üstün,
495 Sanmi Koyejo, Yuntian Deng, Shayne Longpre, Noah A. Smith, Beyza Ermiş, Marzieh
496 Fadaee, and Sara Hooker. The leaderboard illusion. In *Advances in Neural Information*
497 *Processing Systems 38, Datasets and Benchmarks Track*, 2025. arXiv:2504.20879.

- 498 Sandip Sinharay and Paul W. Holland. The missing data assumptions of the NEAT design
499 and their implications for test equating. *Psychometrika*, 75(2):309–327, 2010. doi: 10.
500 1007/s11336-010-9156-6.
- 501 Leon Staufer, Mick Yang, Anka Reuel, and Stephen Casper. Audit cards: Contextualizing
502 ai evaluations, 2025. arXiv:2504.13839.
- 503 Sang Truong, Serena Wang, Nick Haber, and Sanmi Koyejo. Strategic evaluation: Incen-
504 tivizing AI capability coverage with private benchmarks, 2026. Working paper.
- 505 Robert E. Verrecchia. Discretionary disclosure. *Journal of Accounting and Economics*, 5:
506 179–194, 1983. doi: 10.1016/0165-4101(83)90011-3.
- 507 Jialiang Xu, Yifan Mai, and Percy Liang. HELM capabilities: Evaluating LMs capability
508 by capability. Stanford Center for Research on Foundation Models blog, 2025. URL
509 <https://crfm.stanford.edu/2025/03/20/helm-capabilities.html>.
- 510 Chuyifei Zhang, Hongyu Cui, Xiaowen Huang, and Jitao Sang. Positional failures in long-
511 context LLMs: A blind spot in reasoning benchmarks, 2026. arXiv:2605.23170.

512 A What a Provider Could Report

513 The set of results a provider could publish is not fixed, and it grows. When we count the
 514 benchmarks in our panel by the year each was published, the dated stock rises from 11 at
 515 the end of 2019 to 62 by 2026. A benchmark that enters that stock stays available to every
 516 release after it, so the denominator of any omission rate is weakly increasing by construction.

517 What a single release could report grows more slowly and less evenly. Averaged by half-year,
 518 the number of eligible benchmarks per release runs 6.8 in the first half of 2023 and falls to
 519 3.9 in the first half of 2024, a period in which releases outpaced the instruments now used
 520 to score them. By the first half of 2026 it reaches 10.1. The two series move together over
 521 the sample and apart within it, which is consistent with the opportunity set being shaped
 522 partly by release timing rather than by provider choice.

523 A count of omissions therefore measures the calendar. Unless reported tables expand one
 524 for one with the stock, omission rates drift upward for arithmetic reasons alone, and com-
 525 paring early releases with late ones on that measure compares eras rather than providers.
 526 The question that survives is composition: which benchmarks are dropped, conditional on
 527 availability, against what an indifferent selection over the same set would have produced.
 528 Answering it requires knowing not only what was available but how the release stood on
 529 each available benchmark, and it is that second requirement, not the first, that the rest of
 530 the paper shows to be the hard one [Dranove and Jin, 2010].

531 B Additional Checks

532 Is window composition a provider’s choice? Section 3.3 conditions on window composition
 533 and notes that doing so is defensible only if composition is settled before a provider decides
 534 what to report. Regressing the share of a cell’s window released after the focal model on
 535 provider dummies gives an R^2 of 0.071 across 45 organisations, against 0.002 for benchmark
 536 dummies over 61 benchmarks and 0.104 for the release date alone. Window composition is
 537 close to idiosyncratic at the cell level, which is consistent with its being set by the evaluator’s
 538 scoring schedule rather than by the provider, and is not enough to establish it.

539 Is the evaluator’s coverage selective? The estimand conditions on the cells the evaluator
 540 scores, so coverage is upstream of everything else. Coverage averages 8.7 benchmarks per
 541 release with a standard deviation of 8.0 and a range of 1 to 38. It correlates with a provider’s
 542 release count at 0.203 and with the release date at 0.265, and releases from the densest
 543 quartile of providers carry 9.6 scored benchmarks against 8.4 for the rest. Coverage is
 544 therefore mildly tilted toward larger and later providers, and nothing in this paper corrects
 545 for that.

546 Bounding the unidentified component. Write t for a release date, τ for a benchmark’s
 547 vintage and $a = t - \tau$ for the benchmark’s maturity when the model met it. A regression
 548 of standing on all three is singular by construction, which is the age-period-cohort problem,
 549 and the data identify only $\Pi_p = \beta_a + \beta_t$ and $\Pi_c = \beta_\tau - \beta_a$, leaving a one-parameter family
 550 indexed by $\lambda = \beta_a$. Fitting the reduced form on the 2,831 comparable cells gives $\Pi_p = +4.12$
 551 and $\Pi_c = -1.36$ percentile points per year.

552 Three restrictions come from the paper rather than from outside it. Capability rises over
 553 calendar time, so $\beta_t \geq 0$. Later benchmarks are less saturated, which Section 3.2 measures
 554 directly, so $\beta_\tau \leq 0$. Standing rises with maturity, since a model on a young benchmark is
 555 ranked against newer peers, so $\beta_a \geq 0$. Together these give $0 \leq \lambda \leq \min(\Pi_p, -\Pi_c) = 1.36$,
 556 and saturation rather than capability is the binding restriction. The mean maturity gap
 557 between eligible and placebo cells is 2.17 years, so the unidentified component can account
 558 for at most $1.36 \times 2.17 = 2.94$ percentile points, 24% of the observed gap. Adding centred
 559 quadratic and cubic terms in t and τ , which are not collinear with a , moves the ceiling to
 560 34% and 31%.

561 Shuffling scores within benchmark leaves every date and the rank deficiency untouched while
 562 destroying the capability trend. Under that null the same quantity falls to 0.02 percentile

points on average over 30 draws, with a maximum of 0.24, so the bound reflects the trend the geometry multiplies rather than the geometry alone.




Figure 3: The nine HELM Lite models whose published order changed between the first release and the last, ranked among the 24 whose 240 scenario scores are bit-identical throughout; coloured by group of mutually reversing models, four groups and five reversed pairs in all. The control: across 16 HELM Capabilities releases behind an absolute headline, with the pool growing from 22 to 68, none of the 22 constant models’ headline values moved and no pair reversed.

The boundary as a discontinuity. Eligibility is $\mathbf{1}[t_i > \tau_b]$, a deterministic function of the two dates, and that is what the global repairs run into. Equating across groups that took different instruments assumes a conditional relationship common to both, which two groups defined by the boundary cannot have [Sinharay and Holland, 2010], and reweighting by the probability a cell is observed needs that probability bounded away from zero, which a deterministic rule denies [Ma and Chen, 2019]. The same determinism makes the boundary a sharp discontinuity, where standing at the threshold is identified under continuity whatever happens globally [Hahn et al., 2001]. Local linear estimates with a triangular kernel and provider-clustered errors give -8.82 (4.76) at ± 60 days, -6.85 (4.09) at ± 90 , -4.08 (3.11) at ± 180 and -2.85 (2.35) at ± 250 , and a quadratic fit at each width moves none of them across zero. All eight are negative, all eight intervals cover zero, and the largest upper bound is $+8.05$, below the $+12.23$ the global comparison reports.

The local randomization framework agrees [Cattaneo et al., 2015]. Widening symmetric windows and testing the release’s benchmark count, scaffold count and scaffold spread for balance, we find the smallest of the three p -values to be 0.335 at ± 20 days and 0.571 at ± 40 before falling to 0.058 at ± 60 , which selects ± 40 , a rule applied after seeing the balance results, so the randomization p below does not adjust for the selection. The same difference is -5.96 ($p = 0.23$) at ± 20 days and $+0.43$ ($p = 0.88$) at ± 60 . That window holds 113 postdating and 114 eligible cells, and the difference in mean standing is -2.52 with a randomization p of 0.447. Cell counts in the 30 days either side of the cutoff stand at 80 against 84, a ratio of 0.95, so the running variable is not bunching against the threshold [McCrary, 2008].

The mechanism is discontinuous where standing is not. The share of a cell’s window released after the focal model jumps $+0.083$ at the cutoff, which at the eligible-cell gradient of -30.63 implies about -2.5 percentile points, inside every interval above, so at the cutoff the design separates concealment from the availability gap only above that effect, and the narrowest windows cannot reject the full -12.23 either. Peer count and raw score are continuous. The global gap therefore lives in the 793 days that separate the two groups’ mean maturity rather than in anything happening at the boundary itself.

A standing measure with no comparison window. The conclusion points at a scale estimated jointly across benchmarks rather than against contemporaries, which removes the comparison window from the measure entirely. Fitting one does not remove the gap. Ability and difficulty are estimated together over the 63 benchmarks reported on a common zero-to-one scale, on a logit link, with the loading ridged toward one because a free load-

ing is unidentified at this density and takes slopes past 65,000. Restricting to benchmarks carrying at least twenty releases and releases carrying at least three benchmarks leaves 1,669 eligible cells over 31 benchmarks. The model is fit on eligible cells alone, so the 226 postdating cells it can score are out of sample.

The eligible-minus-placebo gap in the residuals is +0.336 residual standard deviations, against +0.449 for the windowed percentile on the same contrast, positive in 64.0% of the 89 releases holding both kinds, at $p = 0.005$. Removing the window shrinks the gap by about a quarter and leaves it positive, which puts a jointly estimated scale beside trimming at 19% and side-balancing at 25%, each on its own scale, rather than ahead of them. Its difficulty parameters are estimated from the maturity-selected cells that exist, so part of the surviving gap may be the same selection re-expressed, and its placebo cells are scored out of sample while every eligible cell is in sample, an asymmetry the contrast inherits. Shuffling scores within benchmark, which leaves every date and the rank deficiency untouched, drops the same quantity to 0.02 on average over 30 draws.




Figure 4: The side-balanced percentile against the windowed one, on eligible cells within release. The windowed measure slopes across the range while the side-balanced one is close to flat where cells are dense; the quoted gradients, -30.63 to -8.56 per unit share, are the within-release regression slopes on these cells, and the release-demeaned binned means shown are correspondingly flatter.

A backward-only standing, position among older peers alone, removes the window's forward half entirely and leaves the label-free contrast at +12.11 across 142 releases, so the gap does not come from look-ahead.

Drop-design ceiling and calibration. The panel contains 144 adjacent same-family release pairs covering 711 shared eligible benchmarks and 15 organisations. Requiring both releases in a pair to carry at least eight scored benchmarks leaves 37 pairs, 482 shared benchmarks, and 8 organisations, the effective cluster count, in the regime where Section 2.5 reports no analytic standard error. Under the null of no selective withholding, the sampling standard deviation of a drop gap, over 400 draws for each of 117 qualifying releases, is 19.4 percentile points when one benchmark is dropped, 14.4 when two are and 12.4 when three are. At conventional two-sided size a single-drop event must therefore move the gap by roughly 38 points before it separates from noise, three times the contamination itself, and no single drop can be read alone.

Possession evidence. Selective withholding is a statement about results a provider held and did not print, and the disclosure label observes only the printing. Four cuts organise how far the record reaches toward possession. The direct cases first: no document explicitly states that a benchmark was evaluated while reporting neither a score nor an outcome, so the direct-withholding category is empty over the 295 cells where a run is documented; the six named-without-number cells all carry an outcome in a figure, a chart, or a comparative claim. Second, the omission deficit stratified by evidence tier. Withholding predicts the deficit weakens, and eventually inverts, as possession evidence strengthens; composition predicts no gradient. The point estimates are +18.9 over 21 releases for dropped benchmarks, +11.3 over 51 for peer-reported omissions, and +11.3 over 69 for silent ones. The deficit does not become more negative as possession evidence strengthens; the dropped tier instead carries the largest positive deficit, a gradient that does not match the signature withholding predicts.

638 Third, a change-based drop test. The estimator above compares levels; a sharper question
639 is whether a benchmark is dropped after the family’s standing on it deteriorates, which
640 differences out persistent salience and relevance. Over the 42 family transitions with a scored
641 predecessor report, 33 dropped and 86 retained cells give a dropped-minus-retained change
642 of -3.7 percentile points, the sign withholding predicts, at a count-preserving permutation
643 $p = 0.16$ within transitions; on raw scores, kept where both sit on the unit interval, the
644 gap is -0.7 points at $p = 0.71$. Both samples cut the 63 derived drop cells: the level
645 estimator keeps the 16 sitting in releases that also retain at least two reported benchmarks
646 to compare against, while the change-based test keeps the 33 with an independent score
647 under both predecessor and successor and imposes no retention floor. Fourth, substitutions.
648 Requiring the reporting table to hold its size while one benchmark is dropped and another
649 inserted, the case a fixed-table explanation cannot cover, leaves no qualifying case in the
650 record. Nothing in the four detects withholding. The change-based gap is directionally
651 consistent with it and statistically indistinguishable from its permutation null; every other
652 cut is null, absent, or opposite in sign.

653 Reliability detail. The second coding covers 20 of the 100 readable releases, drawn strat-
654 ified by provider. Derived categories import contemporaries’ evidence across releases, a
655 dependence the release resampling does not model. Across 15 of them, 194 derived cells
656 pair: one drawn release is a zero-disclosure document whose blank evidence rows the pair-
657 ing drops, a conservative exclusion since all eleven of its cells are trivial agreements, and
658 four drawn releases sit outside the pinned panel. The reported-versus-not cross-tabulation
659 is one-directional, 47 jointly reported and 143 jointly not: no cell the first coder called re-
660 ported was coded otherwise by the second, and the four reversals all sit on one release whose
661 second-coding note records retaining four contested scores despite the first coder’s disagree-
662 ment, an exposure that breaks blindness for that release. Excluding it, reported-versus-not
663 agreement is exact and the full-category κ is 0.88 on 182 cells. Of the 14 disagreements
664 overall, seven are the variant boundary between C and A, one a bare mention against a
665 value, two the derived split between G and H on a single release, and four the exposed
666 release above. Three limits of the record are candidly unrepaired: the eight releases without
667 a locatable document are unlikely to be missing at random and plausibly document less; the
668 evaluator’s coverage may itself respond to what providers emphasise, which would tilt the
669 reported set upward by sampling; and the best-across-scaffolds score rule is not probed for
670 asymmetry across the boundary. High-suspicion marginals match exactly at 5.7% for each
671 coder, which is why the split κ holds at 0.81 despite the category’s rarity; its wide interval
672 reflects eleven such cells concentrating in a few releases.

673 Provider-only clustering, for comparison. The manuscript’s regressions cluster on provider
674 and benchmark jointly. Provider-only errors are consistently smaller: 1.271 against 2.038
675 on the release-effects row, a ratio of 1.60, and 1.652 against 1.748 on the fully saturated
676 row, a ratio of 1.06. The benchmark dimension carries real dependence on the raw contrast
677 and almost none once benchmark effects and the window terms are in, which is what the
678 decomposition says it should. The headline sign-flip test at $p = 0.0001$ preserves provider
679 dependence only; the two-way regression row above is the statement that the contamination
680 survives both dimensions, at $t = -6.6$.

681 NeurIPS Paper Checklist

682 1. Claims

683 Question: Do the main claims made in the abstract and introduction accurately
684 reflect the paper’s contributions and scope?

685 Answer: [\[Yes\]](#)

686 Justification: The abstract and introduction state both findings within their identi-
687 fied scope: the coded omission gap is indistinguishable from the label-free gap, and
688 within-release composition tracks standing. Section 5 reports the reliability of the
689 coding behind them.

690 2. Limitations

691 Question: Does the paper discuss the limitations of the work performed by the
692 authors?

693 Answer: [\[Yes\]](#)

694 Justification: Limitations appear where they bear on the argument rather than
695 only at the end. Section 2.1 bounds the estimand and states that the reported
696 omission rate is descriptive and identifies no withholding; Section 3.2 states that
697 the conditioning is a decomposition rather than a control strategy, and what a
698 study holding disclosure labels would additionally have to argue; Section 3.3 reports
699 the corrections that fail and the coverage the balanced measure loses; Section 5
700 reports the ceiling on the design that narrows endowment uncertainty furthest, its
701 16 realised drops, and the underpowered null they return.

702 3. Theory assumptions and proofs

703 Question: For each theoretical result, does the paper provide the full set of assump-
704 tions and a complete (and correct) proof?

705 Answer: [\[Yes\]](#)

706 Justification: Section 2.1 states one identification result, that a temporal eligibility
707 gate cannot satisfy the exclusion restriction the design needs, and gives its argument
708 in full: the three date variables are exactly collinear, so excluding one requires
709 that at least two of them not move standing. Section 3 shows all three move
710 it. Appendix B states the three sign restrictions used to bound the unidentified
711 component and derives the bound from them. Everything else in Section 2.1 restates
712 published results with citation.

713 4. Experimental result reproducibility

714 Question: Does the paper fully disclose all the information needed to reproduce
715 the main experimental results of the paper to the extent that it affects the main
716 claims and/or conclusions of the paper (regardless of whether the code and data
717 are provided or not)?

718 Answer: [\[Yes\]](#)

719 Justification: Every number in the paper is written by one command into a single
720 machine-readable file, and the data snapshot it was computed against is pinned
721 by a per-file SHA-256 manifest in the released code. Section 2 states the release
722 key, the eligibility rule, the window width, and which estimates impose a minimum
723 benchmark count and which do not.

724 5. Open access to data and code

725 Question: Does the paper provide open access to the data and code, with sufficient
726 instructions to faithfully reproduce the main experimental results, as described in
727 supplemental material?

728 Answer: [\[Yes\]](#)

729 Justification: Code and derived data are released with instructions. The indepen-
730 dent scores are redistributed by their publisher under CC-BY and are fetched rather
731 than copied, then verified against the pinned manifest, so an upstream change is
732 reported rather than silently absorbed. The external leaderboard evidence is frozen

733 with a SHA-256 per source file, because a public leaderboard moves and the argu-
734 ment depends on the exact values.

735 6. Experimental setting/details

736 Question: Does the paper specify all the training and test details (e.g., data splits,
737 hyperparameters, how they were chosen, type of optimizer) necessary to understand
738 the results?

739 Answer: [Yes]

740 Justification: There is no model training. The choices that affect the estimates are
741 stated in Section 2: the release key, the rule for collapsing scaffold variants, the 182-
742 day comparison window, which estimates impose a minimum benchmark count and
743 which do not, and the treatment of pairs whose benchmark vintage is unverified.

744 7. Experiment statistical significance

745 Question: Does the paper report error bars suitably and correctly defined or other
746 appropriate information about the statistical significance of the experiments?

747 Answer: [Yes]

748 Justification: Section 2.5 states the inference protocol. The cluster count is small
749 and unevenly distributed, so below twelve clusters no analytic standard error is re-
750 ported, coefficients in that range use a restricted wild cluster bootstrap, release-level
751 means use a cluster sign-flip randomization test, and no result is called significant
752 on a cluster-t alone. Table 1 reports a standard error with every coefficient, the
753 headline contrast carries a randomization p-value, and the permutation null is re-
754 ported with its mean, its spread and the number of draws reaching the observed
755 value.

756 8. Experiments compute resources

757 Question: For each experiment, does the paper provide sufficient information on the
758 computer resources (type of compute workers, memory, time of execution) needed
759 to reproduce the experiments?

760 Answer: [Yes]

761 Justification: All analysis runs in seconds on a laptop. There is no training, no
762 accelerator use, and no discarded run.

763 9. Code of ethics

764 Question: Does the research conducted in the paper conform, in every respect, with
765 the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

766 Answer: [Yes]

767 Justification: The work uses published benchmark scores and public release docu-
768 ments, involves no human subjects and no personal data, and makes no claim about
769 any individual.

770 10. Broader impacts

771 Question: Does the paper discuss both potential positive societal impacts and neg-
772 ative societal impacts of the work performed?

773 Answer: [Yes]

774 Justification: Section 4 states the consequence for a reader of a pool-relative leader-
775 board, namely that such a number describes standing within one vintage of the eval-
776 uated pool rather than a property of the model. The conclusion states that whether
777 the record needed to measure withholding should grow is a reporting-standard ques-
778 tion. The paper does not attribute concealment to any named provider, because
779 the design does not support that.

780 11. Safeguards

781 Question: Does the paper describe safeguards that have been put in place for re-
782 sponsible release of data or models that have a high risk for misuse (e.g., pre-trained
783 language models, image generators, or scraped datasets)?

784 Answer: [N/A]
785 Justification: No model, dataset or other asset with misuse potential is released.
786 The release is analysis code and derived tables built from already public sources.

787 12. Licenses for existing assets
788 Question: Are the creators or original owners of assets (e.g., code, data, models),
789 used in the paper, properly credited and are the license and terms of use explicitly
790 mentioned and properly respected?
791 Answer: [Yes]
792 Justification: The independent scores are credited to their publisher and used under
793 CC-BY, and the external leaderboard is cited to its published description. Both are
794 named in the text and in the released code.

795 13. New assets
796 Question: Are new assets introduced in the paper well documented and is the
797 documentation provided alongside the assets?
798 Answer: [Yes]
799 Justification: The released code documents how each number is produced, the snap-
800 shot manifest records the exact data build, and the frozen external evidence records
801 a source URL, a retrieval date and a SHA-256 for every file.

802 14. Crowdsourcing and research with human subjects
803 Question: For crowdsourcing experiments and research with human subjects, does
804 the paper include the full text of instructions given to participants and screenshots,
805 if applicable, as well as details about compensation (if any)?
806 Answer: [N/A]
807 Justification: The work involves no crowdsourcing and no human subjects.

808 15. Institutional review board (IRB) approvals or equivalent for research with human
809 subjects
810 Question: Does the paper describe potential risks incurred by study participants,
811 whether such risks were disclosed to the subjects, and whether Institutional Review
812 Board (IRB) approvals (or an equivalent approval/review based on the requirements
813 of your country or institution) were obtained?
814 Answer: [N/A]
815 Justification: The work involves no human subjects, so no review board approval
816 was required.

817 16. Declaration of LLM usage
818 Question: Does the paper describe the usage of LLMs if it is an important, original,
819 or non-standard component of the core methods in this research? Note that if the
820 LLM is used only for writing, editing, or formatting purposes and does not impact
821 the core methodology, scientific rigor, or originality of the research, declaration is
822 not required.
823 Answer: [Yes]
824 Justification: All disclosure-coding categories standing in the paper were hand-
825 assigned by a human coder from archived release-time documents under the frozen
826 protocol, with an independent second coder recoding a provider-stratified 20% sam-
827 ple (Section 5), so no number reported here depends on a model’s judgment. A
828 superseded early pass had used a large language model on the same task; those
829 labels were cleared before analysis and none enters. Every number is produced by
830 the released code from the pinned snapshot. Language models were otherwise used
831 only for writing and code assistance.